

AMAN ARPIT

```
In [116]: # Importing the required Libraries
import pandas as pd
import numpy as np
from datetime import datetime as dt
from scipy import stats

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns; sns.set()

# Data Preprocessing
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder
from sklearn.preprocessing import PowerTransformer

# Model
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.model_selection import cross_val_score, GridSearchCV

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [5]: #Loading the dataset
df = pd.read_csv(r"C:\Users\91700\Downloads\bank+marketing\bank-additional\bank-additional-full")
df_raw = df.copy()
```

```
In [6]: df_raw.head()
```

```
Out[6]:
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	previous	poutcome
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	...	1	999	0	nonexistent
1	57	services	married	high.school	unknown	no	no	telephone	may	mon	...	1	999	0	nonexistent
2	37	services	married	high.school	no	yes	no	telephone	may	mon	...	1	999	0	nonexistent
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	...	1	999	0	nonexistent
4	56	services	married	high.school	no	no	yes	telephone	may	mon	...	1	999	0	nonexistent

5 rows × 21 columns

```
In [7]: df_raw.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    41188 non-null  int64
1   job                    41188 non-null  object
2   marital                41188 non-null  object
3   education              41188 non-null  object
4   default                41188 non-null  object
5   housing                41188 non-null  object
6   loan                   41188 non-null  object
7   contact                41188 non-null  object
8   month                  41188 non-null  object
9   day_of_week            41188 non-null  object
10  duration               41188 non-null  int64
11  campaign               41188 non-null  int64
12  pdays                  41188 non-null  int64
13  previous               41188 non-null  int64
14  poutcome               41188 non-null  object
15  emp.var.rate           41188 non-null  float64
16  cons.price.idx         41188 non-null  float64
17  cons.conf.idx          41188 non-null  float64
18  euribor3m              41188 non-null  float64
19  nr.employed            41188 non-null  float64
20  y                      41188 non-null  object
dtypes: float64(5), int64(5), object(11)
memory usage: 6.6+ MB
```

```
In [9]: df_raw.shape
```

Out[9]: (41188, 21)

In [11]: df_raw.isnull().sum()

Out[11]:

age	0
job	0
marital	0
education	0
default	0
housing	0
loan	0
contact	0
month	0
day_of_week	0
duration	0
campaign	0
pdays	0
previous	0
poutcome	0
emp.var.rate	0
cons.price.idx	0
cons.conf.idx	0
euribor3m	0
nr.employed	0
y	0

dtype: int64

In [12]: df_raw.describe()

Out[12]:

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.e
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	40.02406	258.285010	2.567593	962.475454	0.172963	0.081886	93.575664	-40.502600	3.621291	516.000000
std	10.42125	259.279249	2.770014	186.910907	0.494901	1.570960	0.578840	4.628198	1.734447	7.000000
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	92.201000	-50.800000	0.634000	496.000000
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	93.075000	-42.700000	1.344000	509.000000
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	93.749000	-41.800000	4.857000	519.000000
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	93.994000	-36.400000	4.961000	522.000000
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	94.767000	-26.900000	5.045000	522.000000

Exploratory Data Analysis

In [14]: *#checking for duplicate data*
df_raw.duplicated().sum()

Out[14]: 12

In [19]: *#As there are 12 duplicate values let's drop those values*
print("Before : ",df_raw.shape)
df_n=df_raw.drop_duplicates()
df_n.reset_index(drop=True, inplace=True)
print("After : ",df_n.shape)

print("Number of duplicate data are: ",df_n.duplicated().sum())

Before : (41188, 21)
After : (41176, 21)
Number of duplicate data are: 0

In [20]: *# Preliminary Check*
Variable : age
Description : Customer's age

var ='age'
print('variable :', var)
print()
print('Descriptive stats:')
print(df[var].describe(percentiles=[0.5]))

variable : age

Descriptive stats:
count 41188.00000
mean 40.02406
std 10.42125
min 17.00000
50% 38.00000
max 98.00000
Name: age, dtype: float64

The minimum age is 17, max age is 98 and mean is around 40

The minimum age is 17 ,max age is 98 and mean is around 40

```
In [21]: # Variable : job
# Description : types of job

var='job'
print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())

variable: job

Unique Value Count: 12
['housemaid' 'services' 'admin.' 'blue-collar' 'technician' 'retired'
 'management' 'unemployed' 'self-employed' 'unknown' 'entrepreneur'
 'student']
```

```
In [22]: # Variable : marital
# Description : marital status

var = 'marital'

print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())

variable: marital

Unique Value Count: 4
['married' 'single' 'divorced' 'unknown']
```

```
In [23]: # Variable : education

var = 'education'
print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())

variable: education

Unique Value Count: 8
['basic.4y' 'high.school' 'basic.6y' 'basic.9y' 'professional.course'
 'unknown' 'university.degree' 'illiterate']
```

```
In [24]: # Variable : default

var = 'default'
print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())

variable: default

Unique Value Count: 3
['no' 'unknown' 'yes']
```

```
In [25]: # Variable : housing

var ='housing'

print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())

variable: housing

Unique Value Count: 3
['no' 'yes' 'unknown']
```

```
In [26]: # Preliminary Check
# Variable : month
# Description :last contact month of year
var = 'month'
print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())
print(df[var].dtype)

variable: month

Unique Value Count: 10
['may' 'jun' 'jul' 'aug' 'oct' 'nov' 'dec' 'mar' 'apr' 'sep']
object
```

```
In [27]: #converting months into numerical value
month_rename = {'may':5, 'jun':6, 'jul':7, 'aug':8, 'oct':10, 'nov':11, 'dec':12, 'mar':3, 'apr':4, 'sep':9}
df['month'] = df['month'].map(month_rename).astype(object)

print(df['month'].unique())

[5 6 7 8 10 11 12 3 4 9]
```

```
In [29]: var = 'day_of_week'

print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())
print(df[var].dtype)

variable: day_of_week

Unique Value Count: 5
['mon' 'tue' 'wed' 'thu' 'fri']
object
```

```
In [30]: days_rename = {'mon':1, 'tue':2, 'wed':3, 'thu':4, 'fri':5,}
df['day_of_week'] = df['day_of_week'].map(days_rename)

print(df['day_of_week'].unique())

[1 2 3 4 5]
```

```
In [31]: df.head()
```

```
Out[31]:
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	previous	poutcome
0	56	housemaid	married	basic.4y	no	no	no	telephone	5	1	...	1	999	0	nonexistent
1	57	services	married	high.school	unknown	no	no	telephone	5	1	...	1	999	0	nonexistent
2	37	services	married	high.school	no	yes	no	telephone	5	1	...	1	999	0	nonexistent
3	40	admin.	married	basic.6y	no	no	no	telephone	5	1	...	1	999	0	nonexistent
4	56	services	married	high.school	no	no	yes	telephone	5	1	...	1	999	0	nonexistent

5 rows × 21 columns

```
In [32]: # Description : last contact duration, in seconds (numeric).

var = 'duration'
print('variable:', var)
print()
print('Descriptive stats:')
print(df[var].describe().round())
```

variable: duration

Descriptive stats:

```
count    41188.0
mean       258.0
std        259.0
min         0.0
25%        102.0
50%        180.0
75%        319.0
max        4918.0
Name: duration, dtype: float64
```

As we can see here the difference between min and max value is to high,Here we may have outliers.

```
In [33]: # Description : number of contacts performed during this campaign and for this client

var = 'campaign'

print('variable:', var)
print()
print('Descriptive stats:')
print(df[var].describe().round())
```

variable: campaign

Descriptive stats:

```
count    41188.0
mean         3.0
std          3.0
min          1.0
25%          1.0
50%          2.0
75%          3.0
max         56.0
Name: campaign, dtype: float64
```

```
In [34]: # Description : number of days that passed by after the client was last contacted from a previous campaign

var = 'pdays'

print('variable:', var)
print()
print('Descriptive stats:')
print(df[var].describe().round())

variable: pdays

Descriptive stats:
count      41188.0
mean        962.0
std         187.0
min           0.0
25%         999.0
50%         999.0
75%         999.0
max         999.0
Name: pdays, dtype: float64

we may have outliers here also
```

```
In [35]: # Description : number of contacts performed before this campaign and for this client (numeric)

var = 'previous'

print('variable:', var)
print()
print('Descriptive stats:')
print(df[var].describe().round())

variable: previous

Descriptive stats:
count      41188.0
mean         0.0
std          0.0
min          0.0
25%          0.0
50%          0.0
75%          0.0
max           7.0
Name: previous, dtype: float64
```

```
In [36]: # Variable : poutcome
# Description : outcome of the previous marketing campaign
var = 'poutcome'

print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())
print(df[var].dtype)

variable: poutcome

Unique Value Count: 3
['nonexistent' 'failure' 'success']
object
```

```
In [37]: # Description : employment variation rate

var = 'emp.var.rate'

print('variable:', var)
print()
print('Descriptive stats:')
print(df[var].describe().round())

variable: emp.var.rate
```

```
Descriptive stats:
count      41188.0
mean         0.0
std          2.0
min         -3.0
25%         -2.0
50%          1.0
75%          1.0
max          1.0
Name: emp.var.rate, dtype: float64
```

```
In [38]: df.rename(columns={'emp.var.rate': 'emp_var_rate'}, inplace=True)
```

```
In [39]: # Description : consumer price index - monthly indicator (numeric)
```

```
var = 'cons.price.idx'

print('variable:', var)
print()
print('Descriptive stats:')
print(df[var].describe().round())
```

variable: cons.price.idx

```
Descriptive stats:
count    41188.0
mean       94.0
std         1.0
min        92.0
25%        93.0
50%        94.0
75%        94.0
max        95.0
Name: cons.price.idx, dtype: float64
```

```
In [41]: df.rename(columns={'cons.price.idx':'cons_price_idx'}, inplace=True)
df.rename(columns={'cons.conf.idx':'cons_conf_idx'}, inplace=True)
df.rename(columns={'nr.employed':'nr_employed'}, inplace=True)
```

```
In [42]: # Description : Subscribed or not?
var = 'y'

print('variable:', var)
print()
print('Unique Value Count:', df[var].nunique())
print(df[var].unique())
print(df[var].dtype)
```

variable: y

```
Unique Value Count: 2
['no' 'yes']
object
```

```
In [43]: #let's change it into numeric value that is 0 or 1
y_rename = {'yes':1, 'no':0}
df['y'] = df['y'].map(y_rename)

print(df['y'].unique())
```

[0 1]

```
In [44]: df.rename(columns={'y':'subs_status'}, inplace=True)
```

```
In [48]: df.head(10)
```

```
Out[48]:
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	previous	pou
0	56	housemaid	married	basic.4y	no	no	no	telephone	5	1	...	1	999	0	none
1	57	services	married	high.school	unknown	no	no	telephone	5	1	...	1	999	0	none
2	37	services	married	high.school	no	yes	no	telephone	5	1	...	1	999	0	none
3	40	admin.	married	basic.6y	no	no	no	telephone	5	1	...	1	999	0	none
4	56	services	married	high.school	no	no	yes	telephone	5	1	...	1	999	0	none
5	45	services	married	basic.9y	unknown	no	no	telephone	5	1	...	1	999	0	none
6	59	admin.	married	professional.course	no	no	no	telephone	5	1	...	1	999	0	none
7	41	blue-collar	married	unknown	unknown	no	no	telephone	5	1	...	1	999	0	none
8	24	technician	single	professional.course	no	yes	no	telephone	5	1	...	1	999	0	none
9	25	services	single	high.school	no	yes	no	telephone	5	1	...	1	999	0	none

10 rows × 21 columns

```
In [53]: print(df['subs_status'].value_counts())
print()
print(df['subs_status'].value_counts(normalize=True)*100)
```

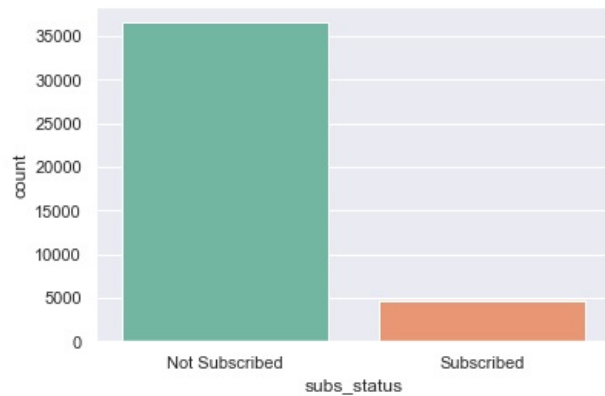
```
0    36548
1     4640
Name: subs_status, dtype: int64
```

```
0    88.734583
1    11.265417
Name: subs_status, dtype: float64
```

```
In [54]: m = sns.countplot(df['subs_status'], palette='Set2')
m.set(xticklabels=['Not Subscribed', 'Subscribed'])
plt.show()
```

C:\Users\91700\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```





```
In [55]: #Analyzing the marital status
fig, ax = plt.subplots(1,2, figsize=(30, 10))
plt.suptitle('Subscription Status by Marital Status')
sns.countplot(df['marital'], hue=df['subs_status'], palette='Set2', order=df['marital'].value_counts().index, ax=ax[0])
ax[0].set_title('Bank Term Deposit count')
graph = df.groupby('marital')['subs_status'].mean().sort_values()
sns.barplot(x=graph.index, y=graph, color='blue', alpha=0.5, ax=ax[1])
ax[1].set_title('Bank Term Deposit Acceptance Rate')

plt.show()
```

C:\Users\91700\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Subscription Status by Marital Status



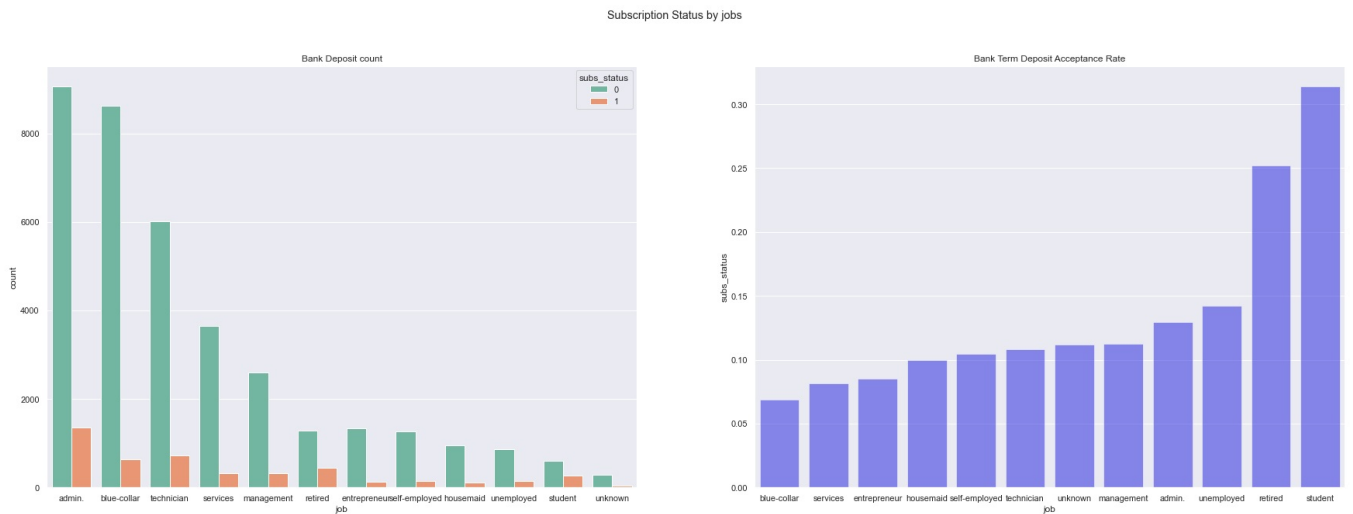
```
In [59]: fig,ax=plt.subplots(1,2, figsize=(30,10))
plt.suptitle('Subscription Status by jobs')
sns.countplot(df['job'], hue=df['subs_status'],palette='Set2', order=df['job'].value_counts().index,ax=ax[0])
ax[0].set_title('Bank Deposit count')

graph = df.groupby('job')['subs_status'].mean().sort_values()
sns.barplot(x=graph.index, y=graph, color='blue', alpha=0.5, ax=ax[1])
ax[1].set_title('Bank Term Deposit Acceptance Rate')

plt.show()
```

C:\Users\91700\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



Student and Retired people have high acceptance rate

Admin and blue collar are frequent clients of the company

```
In [62]: import seaborn as sns
import matplotlib.pyplot as plt

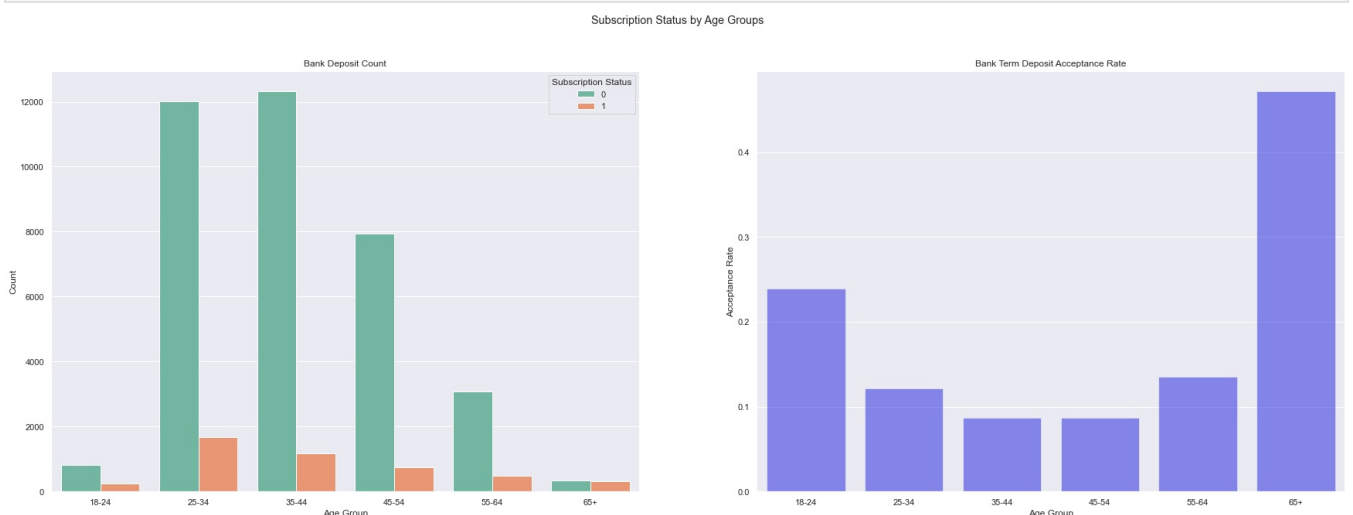
age_groups = ['18-24', '25-34', '35-44', '45-54', '55-64', '65+']
age_bins = [18, 25, 35, 45, 55, 65, 100]
df['age_group'] = pd.cut(df['age'], bins=age_bins, labels=age_groups, right=False)

fig, ax = plt.subplots(1, 2, figsize=(30, 10))
plt.suptitle('Subscription Status by Age Groups')

sns.countplot(x='age_group', hue='subs_status', palette='Set2', order=age_groups, data=df, ax=ax[0])
ax[0].set_title('Bank Deposit Count')
ax[0].set_xlabel('Age Group')
ax[0].set_ylabel('Count')
ax[0].legend(title='Subscription Status')

graph = df.groupby('age_group')['subs_status'].mean().sort_values()
sns.barplot(x=graph.index, y=graph, color='blue', alpha=0.5, ax=ax[1])
ax[1].set_title('Bank Term Deposit Acceptance Rate')
ax[1].set_xlabel('Age Group')
ax[1].set_ylabel('Acceptance Rate')

plt.show()
```



Age group between 25-34 are frequent clients

Age group between 20-34 are frequent clients

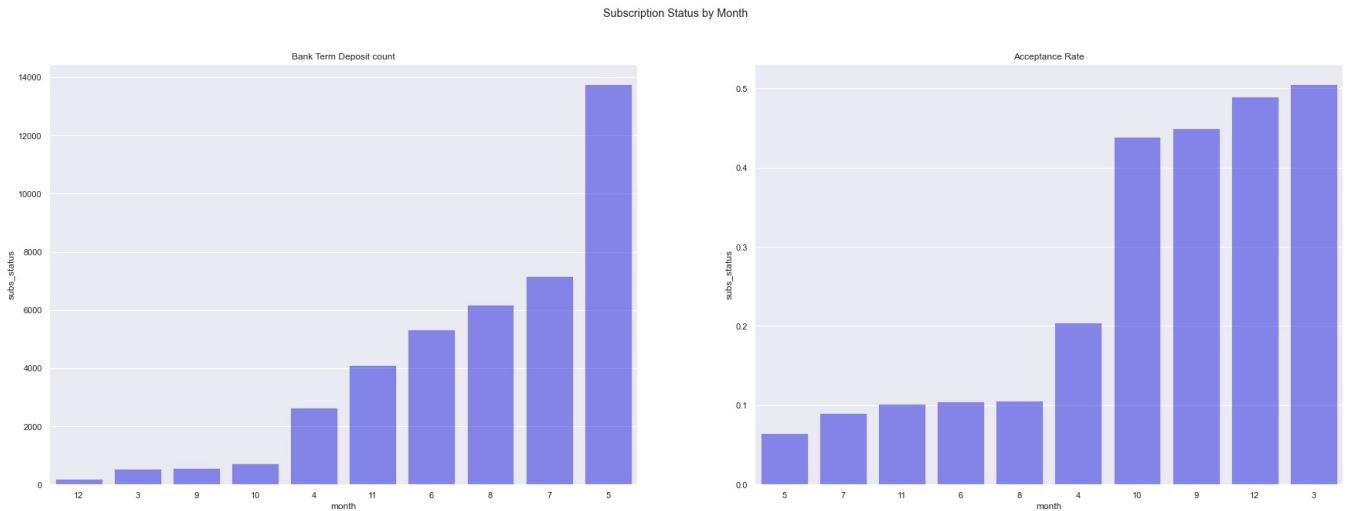
People above the age of 65 have higher acceptance rate

```
In [65]: fig, ax = plt.subplots(1,2, figsize=(30, 10))
plt.suptitle('Subscription Status by Month')

graph = df.groupby('month')['subs_status'].count().sort_values()
sns.barplot(x=graph.index, y=graph, order=graph.index, color='blue', alpha=0.5, ax=ax[0])
ax[0].set_title('Bank Term Deposit count')

graph = df.groupby('month')['subs_status'].mean().sort_values()
sns.barplot(x=graph.index, y=graph, order=graph.index, color='blue', alpha=0.5, ax=ax[1])
ax[1].set_title('Acceptance Rate')

plt.show()
```



In May the clients are highest

.

Data Preprocessing

16 Samples for inference set ,70% Train set and 30% Test set

```
In [66]: random_state = 42
```

```
In [67]: inf_set=df.sample(16,random_state=random_state)
```

```
In [73]: inf_set.shape
```

```
Out[73]: (16, 22)
```

```
In [72]: #Dropping the inference data
```

```
train_test_data=df.drop(inf_set.index)
train_test_data.reset_index(drop=True,inplace=True)
train_test_data.shape
```

```
Out[72]: (41172, 22)
```

```
In [75]: df.shape
```

```
Out[75]: (41188, 22)
```

```
In [77]: # Splitting the data
```

```
# X input parameter only
X = train_test_data.drop(['subs_status'], axis=1)

# y target parameter only
y = train_test_data['subs_status']
```

```
In [106]: X_encoded = pd.get_dummies(X, drop_first=True)
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)
```

C:\Users\91700\Anaconda3\lib\site-packages\pandas\core\algorithms.py:798: FutureWarning: In a future version, the Index constructor will not infer numeric dtypes when passed object-dtype sequences (matching Series behavior)

```

    uniques = Index(uniques)

```

```

In [107.. print("Train set : ",X_train.shape)
          print("Test set : ",X_test.shape)
          print("Inference set : ",inf_set.shape)

```

```

Train set : (32937, 55)
Test set : (8235, 55)
Inference set : (16, 22)

```

```

In [108.. from scipy import stats

z_scores = stats.zscore(df['previous'])
df_no_outliers = df[(z_scores < 3) & (z_scores > -3)]

```

```

In [109.. from scipy import stats

z_scores = stats.zscore(df['duration'])
df_no_outliers = df[(z_scores < 3) & (z_scores > -3)]

```

```

In [110.. from scipy import stats

z_scores = stats.zscore(df['campaign'])
df_no_outliers = df[(z_scores < 3) & (z_scores > -3)]

```

```

In [111.. from scipy import stats

z_scores = stats.zscore(df['pdays'])
df_no_outliers = df[(z_scores < 3) & (z_scores > -3)]

```

```

In [112.. from scipy import stats

z_scores = stats.zscore(df['age'])
df_no_outliers = df[(z_scores < 3) & (z_scores > -3)]

```

```


In [113.. import seaborn as sns
import matplotlib.pyplot as plt

numeric_columns = X_train.select_dtypes(include=['float64', 'int64']).columns

correlation_matrix = X_train[numeric_columns].corr()

plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Correlation Matrix")
plt.show()


```



```

In [153.. df.corrwith(df['subs_status']).sort_values().plot(figsize=(20,10), title='Correlation with `subs_status`',
plt.show()

```



```
In [114]: print('X_train Before', X_train.shape)
print('y_train Before', y_train.shape)
print('X_train After', X_train.shape)
print('y_train After', y_train.shape)
```

```
X_train Before (32937, 55)
y_train Before (32937,)
X_train After (32937, 55)
y_train After (32937,)
```

Model Evaluation

```
In [150]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model = LogisticRegression(random_state=42)

model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(classification_report_str)

y_pred_prob = model.predict_proba(X_test_scaled)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = {:.2f})'.format(roc_auc))
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```


Accuracy: 0.91

Confusion Matrix:

```
[[7080 175]
 [ 555 425]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.98	0.95	7255
1	0.71	0.43	0.54	980
accuracy			0.91	8235
macro avg	0.82	0.70	0.74	8235
weighted avg	0.90	0.91	0.90	8235



```
In [146]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

model = DecisionTreeClassifier(random_state=42)
scaler = StandardScaler()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(classification_report_str)

from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

def plot_roc_curve(model, X, y):
    y_pred_prob = model.predict_proba(X)[:, 1]
    fpr, tpr, _ = roc_curve(y, y_pred_prob)
    roc_auc = auc(fpr, tpr)

    # Plot ROC curve
    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = {:.2f})'.format(roc_auc))
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve')
    plt.legend(loc='lower right')
    plt.show()

plot_roc_curve(model, X_test, y_test)
```


Accuracy: 0.89

Confusion Matrix:

```
[[6803 452]
 [ 495 485]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.94	0.93	7255
1	0.52	0.49	0.51	980
accuracy			0.89	8235
macro avg	0.72	0.72	0.72	8235
weighted avg	0.88	0.89	0.88	8235



```
In [118.. model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(classification_report_str)
```

Accuracy: 0.91

Confusion Matrix:

```
[[7049 206]
 [ 552 428]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.97	0.95	7255
1	0.68	0.44	0.53	980
accuracy			0.91	8235
macro avg	0.80	0.70	0.74	8235
weighted avg	0.90	0.91	0.90	8235

```
In [119.. model = KNeighborsClassifier(n_neighbors=3)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)
```

```

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(classification_report_str)

```

Accuracy: 0.90

Confusion Matrix:

```

[[6900  355]
 [ 502  478]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.95	0.94	7255
1	0.57	0.49	0.53	980
accuracy			0.90	8235
macro avg	0.75	0.72	0.73	8235
weighted avg	0.89	0.90	0.89	8235

In [120..

```

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model = SVC(kernel='rbf', random_state=42)

model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(classification_report_str)

```

Accuracy: 0.90

Confusion Matrix:

```

[[7084  171]
 [ 629  351]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.98	0.95	7255
1	0.67	0.36	0.47	980
accuracy			0.90	8235
macro avg	0.80	0.67	0.71	8235
weighted avg	0.89	0.90	0.89	8235

In [121..

```

model = GaussianNB()

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_report_str = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(classification_report_str)

```

Accuracy: 0.87

Confusion Matrix:

```
[[6665 590]
 [ 508 472]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.92	0.92	7255
1	0.44	0.48	0.46	980
accuracy			0.87	8235
macro avg	0.69	0.70	0.69	8235
weighted avg	0.87	0.87	0.87	8235

In [125...

```
print("LogisticRegression",0.91)
print("DecisionTreeClassifier",0.89)
print("RandomForestClassifier",0.91)
print("KNeighborsClassifier",0.90)
print("GaussianNB",0.87)
print("SVC",0.90)
```

```
LogisticRegression 0.91
DecisionTreeClassifier 0.89
RandomForestClassifier 0.91
KNeighborsClassifier 0.9
GaussianNB 0.87
SVC 0.9
```

Based on the above result , our best performing model is Logistic Regression and The most under performed model in this case is only Naive Bayes model.

Questions

1) What is the distribution of the customer ages?

Ans:- Using this we can easily evaluate

Descriptive stats: count 41188.00000 mean 40.02406 std 10.42125 min 17.00000 50% 38.00000 max 98.00000 Name: age, dtype: float64

The minimum age is 17 ,max age is 98 and mean age is around 40.

2) What is the relationship between customer age and subscription?

Ans:- By analyzing and visualizing the data and graph we can conclude that Age group between 25-34 are frequent clients and People above the age of 65 have higher acceptance rate.

3)Are there any other factors that are correlated with subscription?

Ans:- Based on graph above, we can conclude that:- top three negatively correlated : nr_employed, pdays, euribor3m, emp_var_rate top three positively correlated : duration, and previous

4)What is the accuracy of the logistic regression model?

Ans:- Logistic regression is having the heighest accuracy and it is our best performing model with an accuracy of 0.91.

	precision	recall	f1-score	support
0	0.93	0.98	0.95	7255
1	0.71	0.43	0.54	980
accuracy			0.91	8235

5)What are the most important features for the loaistic rearession

model?

Ans:- a) The logistic regression statistic modeling technique is used when we have a binary outcome variable. For example: will the student pass or fail? Will it rain or not? , will he buy subscription or not? etc.

b) The logistic (or sigmoid) function transforms the log-odds into probabilities between 0 and 1.

c) Logistic regression predicts probabilities rather than class labels

6)What is the precision of the logistic regression model?

Ans:- Precision is a metric used in classification models to evaluate the accuracy of the positive predictions made by the model. It is the ratio of true positive predictions to the total number of positive predictions made by the model.

In our case the precision is around 0.93 for 0 and 0.71 for 1 for Logistic regression model .

7)What is the recall of the logistic regression model?

Ans:- It is used in classification models to evaluate the ability of the model to capture and correctly identify all the relevant instances of the positive class.

In our case recall value is 0.98 for 0 and 0.54 for 1 in logistic regression model.

8)What is the f1-score of the logistic regression model?

Ans:- The F1-score is the harmonic mean of precision and recall, providing a balance between the two metrics. In our case it is 0.95 for 0 and 0.54 for 1.

9)How can you improve the performance of the logistic regression model?

Ans:- Some common method or ideas to improve the performance of the logistic regression model are:-

a) Handle missing data appropriately

b) Encode categorical variables properly (one-hot encoding or label encoding)

c) Identify and handle outliers in the dataset

d) Apply regularization to prevent overfitting (L1, L2 regularization)

e) To Ensure that numerical features are on a similar scale (StandardScaler or MinMaxScaler)

f) Cleaning the data properly and dealing with the missing and duplicate data.

10)What are the limitations of the logistic regression model?

Ans:- a) If the number of observations is lesser than the number of features, Logistic Regression should not be used, otherwise, it may lead to overfitting.

b) The major limitation of Logistic Regression is the assumption of linearity between the dependent variable and the independent variables.

c) It constructs linear boundaries.

d) It is tough to obtain complex relationships using logistic regression. More powerful and compact algorithms such as Neural Networks can easily outperform this algorithm.

END